

Congratulations and Welcome

Udemy Course Structure

1. Snowflake Architecture and Performance Management

- a. Introduction
 - i. Create Trial account
 - ii. Know your UI interface
 - iii. Shared and Non Shared Disk Architecture
 - iv. High Level Architecture
 - v. Create Virtual Warehouses
 - vi. Scaling Policy
 - vii. Create First Database and Table
 - viii. Load CSV from public S3 bucket
- b. Snowflake Architecture
 - i. Snowflake Editions
 - ii. Credits
 - iii. Credit Usage Scenarios
 - iv. Credit to Dollar Storage cost
 - v. Credit Storage Consumption
- c. Performance Management
 - i. Performance Improvement Methods
 - ii. Dedicated Virtual Warehouses
 - iii. Dedicated Virtual Warehouses Creation
 - iv. Scale Up and Scale Down
 - v. Scale Out and Scale In
 - vi. Maximize Caching
 - vii. Maximize Caching demonstration
 - viii. Clustering Keys
 - ix. Clustering Cost
 - x. Clustering Performance
 - xi. Micropartitions

2. Access management and snowflake objects

- a. Access management
 - i. Access control and RBAC
 - ii. Securable objects
 - iii. Roles
- b. Tables in snowflake
 - i. Table Types
 - ii. Database and schema type
- c. User Defined Functions (UDF's) and Stored Procedures
 - i. User defined functions (UDF's)
 - ii. Stored procedures (SP's)
 - iii. Secure UDF's and SP's
- d. Views
 - i. Non Materialized views
 - ii. Materialized views

3. Data Movement - Techniques and Unloading

- a. Data Loading techniques
 - i. Types of Data load
 - ii. Know snowflake stages
 - iii. Install Snowsql
 - iv. Internal Stages
 - 1. Userstage
 - 2. Tablestage
 - 3. Namedstage
 - v. External Stages
 - vi. COPY Command
 - vii. COPY Transformation
 - viii. File Format Object
- b. COPY OPTIONS
 - i. ON ERROR
 - ii. SIZE LIMIT
 - iii. RETURN_FAILED_ONLY
 - iv. ENFORCE_LENGTH / TRUNCATECOLUMNS
 - v. FORCE / LOAD_UNCERTAIN_FILES
 - vi. LOAD_HISTORY
 - vii. VALIDATION MODE
 - viii. Logging rejected records
 - ix. Data loading best practices
- c. Data Unloading
 - i. Data unloading theory
 - ii. COPY into location - Internal Stage -> User stage
 - 1. Resolving max_workers must be greater than 0
 - iii. COPY into location - Internal Stage -> Table stage
 - iv. COPY into location - Internal Stage -> Named Internal Stage
 - v. Unloading using Query
 - vi. Unloading to Multiple Files

4. Data Movement - Data loading

- a. Loading from AWS
 - i. Create AWS free-tier account
 - ii. Create S3 bucket, file and role
 - iii. Storage integration object
 - iv. Load file from AWS
 - v. Unload data to AWS
- b. Loading from Azure
 - i. Create Azure free trial account
 - ii. Create storage and container
 - iii. Integration object
 - iv. Load data from Azure
 - v. Unload data to Azure
- c. Loading from GCP
 - i. Create GCP free trial account, create bucket and upload file
 - ii. Integration object
 - iii. Load data from GCP
 - iv. Unload data to GCP
- d. Loading Semi-Structured Data
 - i. Introduction
 - ii. Process and understanding the demo data
 - iii. JSON: Load raw data and analyze nested objects
 - iv. JSON: FLATTEN
 - 1. Array/Objects
 - 2. Hierarchical
 - v. JSON: Data transformation and load
 - vi. JSON: Querying Data
 - vii. PARQUET: File load and transformation
- e. Snowpipe
 - i. Overview
 - ii. Create snowpipe, load and configure
 - iii. Troubleshooting and pipe refresh
 - iv. Manage snowpipe

5. Dynamic Data Masking and Visualization

- a. Dynamic data masking
 - i. Overview
 - ii. Create masking policy
 - iii. Change, Unset and replace masking policy
 - iv. Conditional masking

- b. Data Visualization
 - i. Snowsight - Dashboard and Visualization
 - ii. Getting started with Tableau
 - iii. Tableau visualization from snowflake data

- c. Partner Connect
 - i. Overview
 - ii. Partner connect categories

6. Data Sharing and Data Protection

a. Data Sharing

- i. Overview
- ii. Data sharing demo of Table, External tables and UDF's.
- iii. Data share to a reader account
- iv. Imported privileges
- v. Data share - entire database
- vi. Data share operations - DDL and DML affect
- vii. Sharing views and materialized views

b. Time Travel

- i. Overview and retrieve data using "at" and "before"
- ii. Data restoration using direct method
- iii. Data restoration using indirect method
- iv. UNDROP
- v. Understanding retention time

c. Fail-Safe

- i. Overview
- ii. Fail Safe demo

d. Zero-Copy Cloning

- i. Overview
- ii. Cloning permanent tables
- iii. Copy grants and clone temporary external tables
- iv. Clone schema and database
- v. Swap
- vi. Clone using time travel and understand sharing vs cloning

e. Streams

- i. Overview
- ii. Consuming streams using insert, update and delete
- iii. Types of streams - default, append only. Consume stream using false condition
- iv. Insert only streams for external tables
- v. Multiple consumers with streams

- vi. Changes clause
- vii. Streams consideration and stale after attribute
- viii. Recreate, rename and clone effect on stream

f. Tasks

- i. Understanding tasks
- ii. Create task using procedure when and troubleshoot task
- iii. Scheduling tasks using CRON
- iv. CRON - Daylight savings considerations
- v. DAG - Tree of tasks
- vi. Task with stream
- vii. User managed tasks and warehouse considerations

g. Sampling Data

- i. Overview
- ii. Row sampling method
- iii. System block method

h. Access management in practice

- i. System defined roles and its users
- ii. Custom roles and its users
- iii. Create objects from sysadmin role and assign to developer role
- iv. Create custom users - developer1, developer2, qa1 and reader1
- v. Problems due to unassigned custom roles
- vi. Solution of the problems
- vii. Public Role

i. Best Practices

- i. Auto resume, success and failure and long running queries
- ii. Resource monitor
- iii. Columns storage cost and benchmarking the queries